

COMP/EECE 7/8740 Neural Networks

Topics:

Logistic regression

- Logistic regression
- Underfitting or overfitting problems and
- Solutions

Md Zahangir Alom
Department of Computer Science
University of Memphis, TN

What is Logistic Regression ?

- This type of statistical model is often used for **classification and predictive analytics**.
- Logistic regression **estimates the probability of an event occurring**, such as voted or didn't vote, based on a given dataset of independent variables.
 - Since the outcome is a probability, **the dependent variable is bounded between 0 and 1**.
- Similar to linear regression, logistic regression is also used to estimate the relationship between a dependent variable and one or more independent variables, but it is used to **make a prediction about a categorical variable versus a continuous one**.
 - A categorical variable can be true or false, yes or no, 1 or 0, et cetera.

Linear versus Logistic Regression

- The unit of measure also differs from linear regression as it produces a probability, **but the logit function transforms the S-curve into straight line**
- Both models are used in regression analysis to make predictions about future outcomes, **linear regression is typically easier to understand**
- Linear regression also does not require as large of a sample size as **logistic regression needs an adequate sample** to represent values across all the response categories.
- Without a larger, representative samples, the model may not have sufficient statistical power to detect a significant effect.

Types of logistic regression

- There are **three types of logistic regression models**, which are defined based on categorical response.
 - **Binary logistic regression:** In this approach, the response or dependent variable is dichotomous in nature—i.e. it **has only two possible outcomes** (e.g. 0 or 1). Some popular examples of its use include predicting if an e-mail is spam or not spam or if a tumor is malignant or not malignant (most popular type).
 - **Multinomial logistic regression:** In this type of logistic regression model, the dependent variable **has three or more possible outcomes**; however, these values have no specified order. For example, a multinomial logistic regression model can help the studio to determine the strength of influence a person's age, gender, and dating status may have on the type of film that they prefer.
 - **Ordinal logistic regression:** This type of logistic regression model is leveraged when the response variable has **three or more possible outcome, but in this case, these values do have a defined order**. Examples of ordinal responses include grading scales from A to F or rating scales from 1 to 5.


Logistic regression in machine learning

- In Machine Learning (ML), logistic regression belongs to the family of supervised learning models.
- The **negative log likelihood used as the loss function**, using the process of gradient descent to find the global maximum.
- Logistic regression can also **be prone to overfitting**, particularly when there is a high number of predictor variables within the model.
- **Regularization is typically used to penalize parameters large coefficients** when the model suffers from high dimensionality.

CLASSIFICATION


Classification

- Target value – 0 or 1 – 2 Classes




Some Examples

- Email: Spam / Not Spam
- Online Transaction: Fraudulent (Yes/No)
- Tumor: No/Yes



No



Yes

Target Variable 'Y'

- For 2-class problem

– $y \in \{0,1\}$

0 : Negative Class : No Tumor

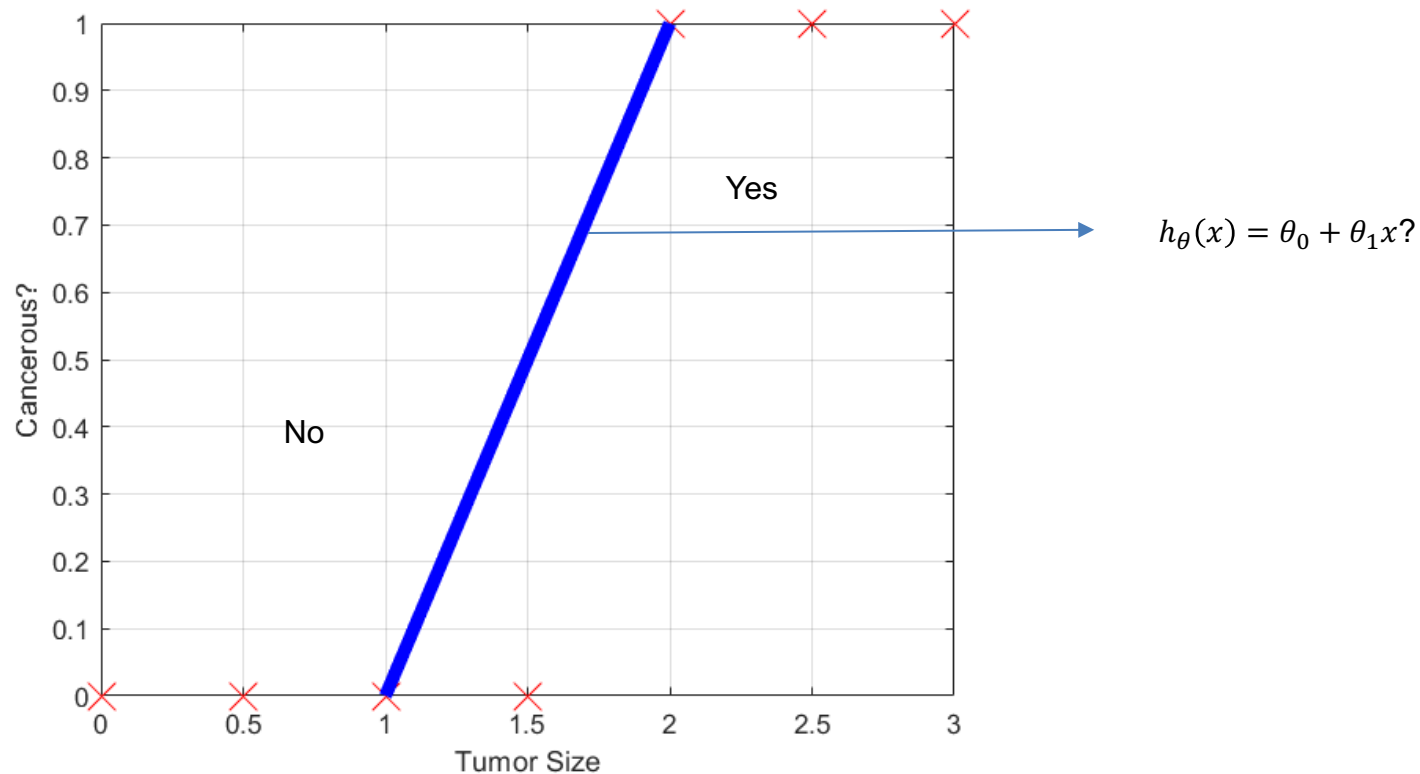
1 : Positive Class : Tumor

- For 4-class problem

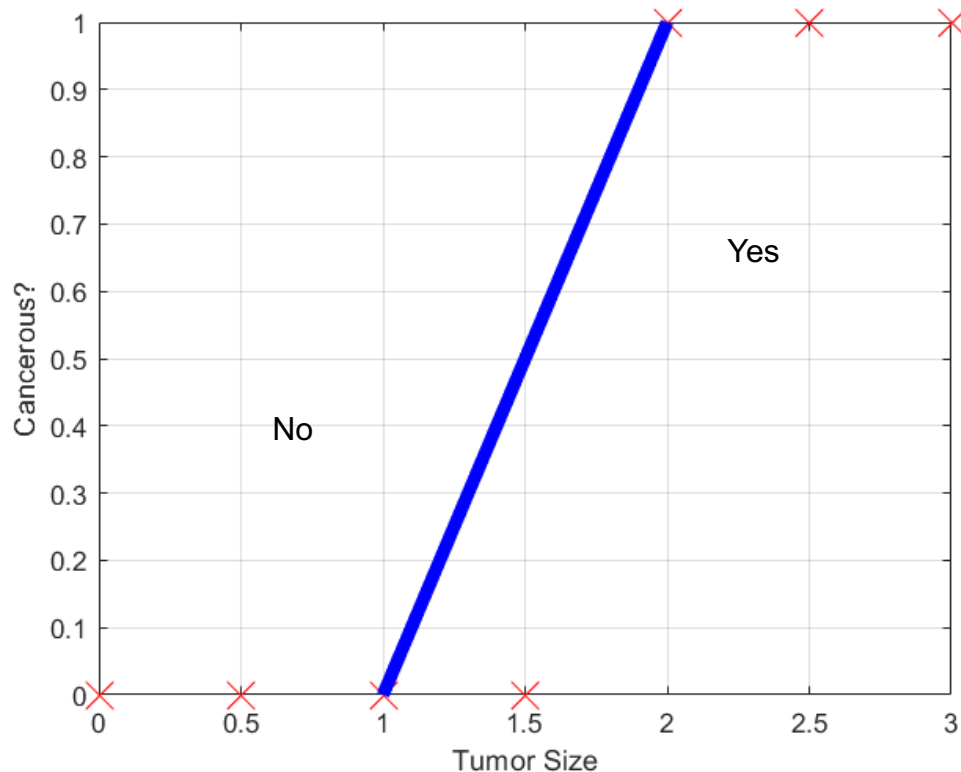
– $y \in \{0,1,2,3\}$



Let's start with 2 class variable




Conversion of Regression to Classification?



$h_{\theta}(x) = x - 1$
if $h_{\theta}(x) \geq 0.5$ \longrightarrow Predict $y(\hat{y})=1$
 $h_{\theta}(x) < 0.5$ \longrightarrow Predict $y(\hat{y})=0$

Let's add one more training example?



Poor performance in this scenario

“Adding one example has reduced the performance considerably”

“Other funny things can happen too”


‘h’ should be technically 0 or 1 but there is a possibility that ‘h’ can be a greater value as $\theta_0 + \theta_1 x$

Logistic regression

- Don't be confused by the name: It's designed for **classification**
- Classification Algorithm: $0 \leq h_{\theta}(x) \leq 1$
- Hypothesis:
 - $h_{\theta}(x) = \theta x^T$ - Linear Regression
 - $h_{\theta}(x) = g(\theta x^T)$ - Logistic Regression
 - g ?

We have to modify the linear regression equation slightly to find the solution for logistic regression

Sigmoid function




$$g(z) = \frac{1}{1 + e^{-z}}$$

Interpretation of Hypothesis

- $h_{\theta}(x)$: estimated probability that $y = 1$ for the input 'x'
- Let's say, $x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ Tumor\ Size \end{bmatrix}$ and let's say our algorithm predicts $h_{\theta}(x) = 0.7$, then 70% of chance of tumor being cancerous
- $h_{\theta}(x) = P(y = 1 \mid x; \theta) \rightarrow$ Probability that $y=1$ given 'x' estimated by θ
- $P(y = 0 \mid x; \theta) = 1 - P(y = 1 \mid x; \theta)$

Decision Boundary


- $h_{\theta}(x) = g(\theta x^T) = P(y = 1 \mid x; \theta)$
- $g(z) = \frac{1}{1+e^{-z}}$
- $\hat{y} = 1$ if $h_{\theta}(x) \geq 0.5$
- $\hat{y} = 0$ if $h_{\theta}(x) < 0.5$
- $g(z) \geq 0.5$ whenever $z \geq 0$
- $g(\theta x^T) \geq 0.5$ whenever $\theta x^T \geq 0$




Predicts $y = 1$ whenever $\theta x^T \geq 0$

Predicts $y = 0$ whenever $\theta x^T < 0$

When the Hypothesis value is 0.5



Decision Boundary




$$\theta(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

$$\theta_0 = -3, \theta_1 = 1, \theta_2 = 1$$


- $\hat{y} = 1$ if $-3 + x_1 + x_2 \geq 0$
- $\hat{y} = 0$ if $-3 + x_1 + x_2 < 0$
- If $x_1 + x_2 \geq 3 \rightarrow \hat{y} = 1$
- If $x_1 + x_2 < 3 \rightarrow \hat{y} = 0$

Non-linear Decision Boundary



How can we make this happen?

Non-linear Decision Boundary



Hypothesis:

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

- $\theta_0 = -1, \theta_1 = 0, \theta_2 = 0, \theta_3 = 1, \theta_4 = 1$

We can add higher order polynomial features as needed

$$x_1^2 + x_2^2 \geq 1 \rightarrow \hat{y} = 1$$

$$x_1^2 + x_2^2 < 1 \rightarrow \hat{y} = 0$$

Logistic regression

- Logistic regression can be used to **determine very complex boundaries** if necessary
- Problem definition
 - Training Set: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}) \dots (x^{(m)}, y^{(m)})\}$
 - $x \in [x_0, x_1, x_2, x_3 \dots x_n]$
 - $y \in [0,1]$
 - $$h_{\theta}(x) = \frac{1}{1+e^{-\theta x^T}}$$
- How to choose θ ?

Linear regression cost function

- $J = \frac{1}{2m} \sum_{i=1}^M (h_{\theta}(x^i) - y^i)^2$ - from previous class

Can this be applied for logistic regression?

- Answer is yes and no (depends)
- Mean square error performs **reasonably well but not consistent**
- Mean square error are used for certain classification problems, however, **there is a possibility of local minima** ,
Not a convex function always!
- Gradient descent for classification using **mean square error(MSE) might not converge to global minima**

Logistic regression cost function

$$P(y | x) = \begin{cases} \hat{p} & \text{if } y = 1 \\ 1 - \hat{p} & \text{if } y = 0 \end{cases}$$

Max/Min occurs at points where derivative is equal to 0:

$$\frac{\partial P(y | x)}{\partial p} = 0$$

Our examples can only be part of one class at a time either 1 or 0 so,

- when $y = 1$ then \hat{p} should be the output and $1 - \hat{p}$ should be ignored,
- similarly, when $y = 0$ then $1 - \hat{p}$ should be the output and \hat{p} should be ignored.

We can combine this as follows:

$$P(y | x) = \hat{p}^y (1 - \hat{p})^{1-y}$$

if $y = 1$ then:

$$P(y | x) = \hat{p}^y (1 - \hat{p})^{1-y} = \hat{p}^1 (1 - \hat{p})^{1-1} = \hat{p}$$

if $y = 0$ then:



$$P(y | x) = \hat{p}^y (1 - \hat{p})^{1-y} = \hat{p}^0 (1 - \hat{p})^{1-0} = 1 - \hat{p}$$

Logistic regression cost function

Recall the properties of natural log:



- $\log(a * b) = \log(a) + \log(b)$
- $\log(a^x) = x \log(a)$

$$\begin{aligned}\log(P(y | x)) &= \log(\hat{p}^y (1 - \hat{p})^{(1-y)}) \\ &= \log(\hat{p}^y) + \log((1 - \hat{p})^{(1-y)}) \\ &= y \log(\hat{p}) + (1 - y) \log(1 - \hat{p})\end{aligned}$$




Logistic regression cost function

$$\begin{aligned}\text{Binary Cross Entropy Loss}(y, \hat{p}) &= -\log(P(y | x)) \\ &= -(y \log(\hat{p}) + (1 - y) \log(1 - \hat{p})) \\ &= -y \log(\hat{p}) - (1 - y) \log(1 - \hat{p})\end{aligned}$$



Logistic regression cost function



Logistic regression cost function

$$\mathcal{L}(\hat{p}^{(1)}, \hat{p}^{(2)}, \dots, \hat{p}^{(m)} \mid (x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})) = \prod_{i=1}^m \hat{p}^{(i)^{y^{(i)}}} (1 - \hat{p}^{(i)})^{1-y^{(i)}}$$

This is the **likelihood**(\mathcal{L}) function, where we are trying to maximize the each probability, \hat{p} , given the examples $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})$.

$\prod_{i=1}^m$ represents product over examples 1 to m

$$\begin{aligned} \log \left(\mathcal{L}(\hat{p}^{(1)}, \hat{p}^{(2)}, \dots, \hat{p}^{(m)} \mid (x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})) \right) &= \log \left(\prod_{i=1}^m \hat{p}^{(i)^{y^{(i)}}} (1 - \hat{p}^{(i)})^{1-y^{(i)}} \right) \\ &= \sum_{i=1}^m \log \left(\hat{p}^{(i)^{y^{(i)}}} (1 - \hat{p}^{(i)})^{1-y^{(i)}} \right) \\ &= \sum_{i=1}^m y^{(i)} \log(\hat{p}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{p}^{(i)}) \end{aligned}$$

This is called the **log-likelihood** function

Logistic regression cost function

$$\begin{aligned}\text{Binary Cross-Entropy Cost} &= \frac{1}{m} \left(-\log \left(\mathcal{L}(\hat{p}^{(1)}, \hat{p}^{(2)}, \dots, \hat{p}^{(m)}) \right) \right) \\ &= \frac{1}{m} \sum_{i=1}^m -y^{(i)} \log(\hat{p}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{p}^{(i)})\end{aligned}$$

$$\begin{aligned}-\log \left(\mathcal{L}(\hat{p}^{(1)}, \hat{p}^{(2)}, \dots, \hat{p}^{(m)}) \right) &= -\sum_{i=1}^m y^{(i)} \log(\hat{p}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{p}^{(i)}) \\ &= \sum_{i=1}^m -y^{(i)} \log(\hat{p}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{p}^{(i)})\end{aligned}$$

Logistic regression cost function

- $Cost(h_{\theta}(x^i) - y^i) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$

Since y is either going to be 0 or 1

- $J = \frac{1}{m} [\sum_{i=1}^M -y^i \log(h_{\theta}(x^i)) - (1 - y^i) \log(1 - h_{\theta}(x^i))]$

Principle of Maximum Likelihood

- $J = \frac{-1}{m} [\sum_{i=1}^M y^i \log(h_{\theta}(x^i)) + (1 - y^i) \log(1 - h_{\theta}(x^i))]$

Gradient Descent(GD)

- To fit parameters θ
- To make a prediction :

- $h_{\theta}(x) = \frac{1}{1+e^{-\theta x^T}}$

- $J = \frac{-1}{m} [\sum_{i=1}^M y^i \log(h_{\theta}(x^i)) + (1 - y^i) \log(1 - h_{\theta}(x^i))]$
- Minimize $J(\theta)$

Gradient Descent – Update θ

- $\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$ - Similarly update all θ_j
- $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{-1}{m} \frac{\partial}{\partial \theta_j} [\sum_{i=1}^M y^i \log(h_\theta(x^i)) + (1 - y^i) \log(1 - h_\theta(x^i))]$

Gradient Descent Equation

- $$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^M (h_{\theta}(x^i) - y^i) x^i$$


Exactly the same as Linear Regression except the
$$h_{\theta}(x^i)$$

Logistic Regression Steps


- Load the dataset: Input Features and Target Variable
- Normalize Features
 - $\hat{x} = \frac{x - \mu}{\sigma}$ or any other methods
 - Add $x_0 = 1$ to all samples (Normalized \hat{x})
 - new Dimensions of x – $M \times (n+1)$
- Implement cost function after computing hypothesis
 - $h_{\theta}(x) = g(\theta x^T), g(z) = \frac{1}{1+e^{-z}}$
 - $J = \frac{-1}{m} [\sum_{i=1}^M y^i \log(h_{\theta}(x^i)) + (1 - y^i) \log(1 - h_{\theta}(x^i))]$
 - No need of for loops in this scenario too
- Implement gradient descent function
 - $\theta = \theta - \alpha \delta$
 - $\delta = \frac{1}{m} \sum_{i=1}^M (h_{\theta}(x^i) - y^i) x^i$
- Compute $h_{\theta}(x)$ with updated θ
- Threshold $h_{\theta}(x)$ with 0.5 and determine $\hat{y} \in \{0,1\}$
- Evaluate the performance

Logistic regression for multi-class variable


- Email Folder: Work, Personal, Spam, Ads



- Medical Imaging: No Cancer, Benign, Malignant




- Weather: Sunny, Cloudy, Rain, Snow








Binary vs. Multi-Class Classification

Binary classification:



Multi-class classification:



-  • Class 1:
-  • Class 2: 
-  • Class 3:

Multi-Class Classification

- 3 Binary classification approaches:
 - Where class 1 is positive, classes 2 & 3 are negative
 - Where class 2 is positive, classes 1 & 3 are negative
 - Where class 3 is positive, classes 1 & 2 are negative

One vs. All

- Train a logistic regression classifier $h_{\theta}(x^i)$ for each class ' i ' to predict the probability that $y = i$
- On a new input x , to make a prediction, pick the class ' i ' that maximize $\Rightarrow \max_i (h_{\theta}^i(x))$
 - In this case, '3' classifiers – Pick the one with highest probability

PROBLEM OF UNDERFITTING /OVERFITTING

Regularization with **Ridge, Lasso, and Elastic Net Regressions**


- ***What is Regularization:*** Techniques for combating overfitting and improving training.
- A common phenomenon referred to as “overfitting”. Overfitting has **a polar opposite called underfitting**
- In technical terms, overfitting means the model you built **has more parameters than the data can justify**
- You may be asking *what does all this have to do with regularization?*
 - The answer is **everything**.

Measure to address overfitting

- Reduce number of features
 - Choose manually
 - Feature selection algorithm (Later in the course)
- Regularization
 - Keep all features, but **reduce magnitude/values of θ_j**
 - When we have a lot of features, each feature can make an important contribution to predict 'y'
 - This method would help us preserve all features
- Overview of the differences in 3 common regularization techniques:
 - Ridge
 - Lasso, and
 - Elastic Net.

Regularization with Ridge, Lasso, and Elastic Net Regressions


- **Adding the right amount of bias to a model can help make more accurate predictions by desensitizing it to some of the noise in the training data.**



In the plots above the **blue dots** are sample data points taken from the real-world. The distance those samples are from the **yellow line**, “True Function”, is called the “noise” of the **data**. The distance of the sample points to the blue line is referred to as the “error” of our model.

Ridge regression : regularized linear regression cost function

- $$J = \frac{1}{2m} \left[\sum_{i=1}^M (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$



Shrink all parameters - reduces over reliance on any feature

- By convention, regularization is done only for θ_1 to θ_n and not for θ_0
- Small values for parameters:
 - “Simpler” Hypothesis
 - Less prone to overfitting

Ridge regression : regularized linear regression cost function

- $$J = \frac{1}{2m} \left[\sum_{i=1}^M (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

Keeping the parameters small to avoid overfitting

Fit the training set

Regularization
Parameter: Trade off

- λ is adjusted to control the value of θ and minimal error
- Let's say, $\lambda = 10^3$, $\theta_j \cong 0$, $j = 1, 2, \dots, n$ which would mean $h_{\theta}(x) = \theta_0$ and end up in underfitting
- λ should be chosen carefully

Ridge regression : regularized linear regression

- Regularized Linear Regression

- $J = \frac{1}{2m} \left[\sum_{i=1}^M (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$

- Gradient descent without regularization

- $\theta = \theta - \alpha \delta$

- $\delta = \frac{1}{m} \sum_{i=1}^M (h_{\theta}(x^i) - y^i) x^i$

$$h_{\theta}(x) = \theta x^T$$

- Gradient descent with regularization $x_0 = 1$

- $\theta_0 = \theta_0 - \frac{\alpha}{m} \sum_{i=1}^M (h_{\theta}(x^i) - y^i) x_0^i$

- $\theta_j = \theta_j - \frac{\alpha}{m} \left[\sum_{i=1}^M (h_{\theta}(x^i) - y^i) x_j^i + \lambda \theta_j \right]$

- $\theta_j = \theta_j \left[1 - \frac{\alpha \lambda}{m} \right] - \frac{\alpha}{m} \left[\sum_{i=1}^M (h_{\theta}(x^i) - y^i) x_j^i \right]$

Ridge regression : regularized logistic regression

- Regularized Logistic Regression

- $$J = \frac{-1}{m} [\sum_{i=1}^M y^i \log(h_{\theta}(x^i)) + (1 - y^i) \log(1 - h_{\theta}(x^i))] + [\frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2]$$

- Gradient descent without regularization

- $$\theta = \theta - \alpha \delta$$

- $$\delta = \frac{1}{m} \sum_{i=1}^M (h_{\theta}(x^i) - y^i) x^i$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta x^T}}$$

- Gradient descent with regularization

$$x_0 = 1$$

- $$\theta_0 = \theta_0 - \frac{\alpha}{m} \sum_{i=1}^M (h_{\theta}(x^i) - y^i) x_0^i$$

- $$\theta_j = \theta_j - \frac{\alpha}{m} [\sum_{i=1}^M (h_{\theta}(x^i) - y^i) x_j^i + \lambda \theta_j]$$

- $$\theta_j = \theta_j \left[1 - \frac{\alpha \lambda}{m}\right] - \frac{\alpha}{m} [\sum_{i=1}^M (h_{\theta}(x^i) - y^i) x_j^i]$$

Ridge regression : regularization technique

- This function **penalizes your model for having too many or too large predictors**
- The objective of Ridge regression is to reduce the effect of these predictors to **decrease the chance of overfitting your data.**
- If we were to set $\lambda = 0$ then this would be a normal linear regression.
- The most common use of Ridge regression is to be *preemptive* in **addressing overfitting concerns**
- Ridge regression is a good tool for handling multicollinearity *when you must keep all your predictors.*

Ridge regression works well if there are many predictors of about the same magnitude. This means all predictors have similar power to predict the target value.

Ridge regression : reference codes

```
from sklearn.model_selection import Ridge

# Create an array of alpha values to test
alphas = np.logspace(-1, 1.5, 500, base=10)

# Create a Ridge regression model instance
ridge = Ridge(random_state=0, max_iter=10000, alpha=alphas)

# Create dictionary key,value pair of alpha values
tuned_parameters = [{'alpha': alphas}]

# Specify number of folds for cross_validation
n_folds = 5

# Create grid search instance using desired variables
clf_ridge = GridSearchCV(ridge, tuned_parameters, cv=5, refit=False)
clf_ridge.fit(x_train_scaled, y_train)
ridge_scores = clf_ridge.cv_results_['mean_test_score']

# Plot the Figure
plt.figure().set_size_inches(8, 6)
plt.plot(alphas, ridge_scores)
plt.xlabel('Alpha Value')
plt.ylabel('Cross Validation Score')
plt.title('Ridge Regression Alpha Demonstration')
plt.axvline(clf_ridge.best_params_['alpha'], color='black',
            linestyle='--')
print(f'The optimal alpha value is:
{clf_ridge.best_params_["alpha"]}')

```

Lasso Regression

: regularization techniques

- **Lasso regression** the equation below and thinking to yourself “that looks almost identical to Ridge regression.”.
- Lasso differs from Ridge regression by summing the **absolute value of the predictors (m_j)** instead of summing the squared values.

$$\text{cost_function_lasso} = \sum_{i=1}^n (y_i - \sum_{j=1}^k (m_j x_{ij}) - b)^2 + \lambda \sum_{j=1}^p |m_j|$$

Lasso Cost Function — Notice the lambda (λ) multiplied by the sum of the absolute value of the predictors

- Lasso is an acronym that stands for “Least Absolute Shrinkage and Selection Operator.” ***Due to the penalty term not being squared, some values can reach 0. When a predictor coefficient (m_j) reaches 0 that predictor does not affect the model.***
- Lasso tends to do well **if there are few significant predictors and the** magnitudes of the others are close to zero.

Lasso regression : reference codes

```
from sklearn.model_selection import GridSearchCV

# Create an array of alpha values to test
# Start np.linspace value is 10**-10 because a value of 0 throws
warnings
alphas = np.logspace(-10, 1, 1000, base=10)

# Create dictionary key,value pair of alpha values
tuned_parameters = [{'alpha': alphas}]

# Specify number of folds for cross_validation
n_folds = 5

# Create grid search instance using desired variables
clf_lasso = GridSearchCV(lasso, tuned_parameters, cv=5, refit=True)
clf_lasso.fit(x_train_scaled, y_train)
lasso_scores = clf_lasso.cv_results_['mean_test_score']

# Plot the results
plt.figure().set_size_inches(8, 6)
plt.plot(alphas, lasso_scores)
plt.xlabel('Alpha Value')
plt.ylabel('Model CV Score')
plt.title('Lasso Regression Alpha Demonstration')
plt.axvline(clf_lasso.best_params_['alpha'], color='black',
            linestyle='--')
print(f'The optimal alpha value is :
{clf_lasso.best_params_["alpha"]}')

```


Elastic Net Regression

: regularization technique

- **Elastic Net regression** was created as a critique of Lasso regression.
- **Elastic Net** is a combination of both Lasso and Ridge regressions. .

$$\operatorname{argmin}_{\beta} \sum_i (y_i - \beta' x_i)^2 + \lambda_1 \sum_{k=1}^K |\beta_k| + \lambda_2 \sum_{k=1}^K \beta_k^2$$

The λ_1 is Lasso penalty (L1) and λ_2 is the Ridge regression penalty (L2)

Elastic Net Regression

: regularization technique

- As you can see in the picture above there are now two λ terms. λ_1 is the “alpha” value for the Lasso part of the regression and λ_2 is the “alpha” value for the Ridge regression equation.
- When using sci-kit learn’s Elastic Net regression the alpha term is a ratio of $\lambda_1:\lambda_2$
- Setting the ratio values:
 - **ratio = 0 it acts as a Ridge regression, and**
 - **when the ratio = 1 it acts as a Lasso regression.**
 - **Any value between 0 and 1 is a combination of Ridge and Lasso regression**

Summary : regularization technique

- ***When do I use Regularization:***

- Ridge regression — Ridge regression works well if there are many large predictors of about the same value.
- Lasso — Few significant predictors and the magnitudes of the others are close to zero
- Elastic Net — Mixture of both Ridge and Lasso

- ***How do I use Regularization:***

- Split and Standardize the data (only standardize the model inputs and not the output)
- Decide which regression technique Ridge, Lasso, or Elastic Net you wish to perform.
- Use **GridSearchCV** to optimize the hyper-parameter alpha
- Build your model with optimized alpha and make predictions!

Summary

- Logistic regression
- Overfitting and under fitting problems
- Solutions with regularization techniques
- What's next ?
 - Support Vector Machine (SVM)
 - Kernelization approaches